

ex1

June 13, 2021

```
[1]: '''Trains a simple deep NN on the MNIST dataset.
```

```
Gets to 98.40% test accuracy after 20 epochs  
(there is *a lot* of margin for parameter tuning).  
2 seconds per epoch on a K520 GPU.  
'''
```

```
[1]: 'Trains a simple deep NN on the MNIST dataset.\n\nGets to 98.40% test accuracy  
after 20 epochs\n(there is *a lot* of margin for parameter tuning).\n2 seconds  
per epoch on a K520 GPU.\n'
```

```
[2]: from tensorflow import keras  
from tensorflow.keras.datasets import mnist  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Dropout  
from tensorflow.keras.optimizers import RMSprop
```

```
[3]: batch_size = 128  
num_classes = 10  
epochs = 20
```

```
[4]: # the data, split between train and test sets  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
  
x_train = x_train.reshape(60000, 784)  
x_test = x_test.reshape(10000, 784)  
x_train = x_train.astype('float32')  
x_test = x_test.astype('float32')  
x_train /= 255  
x_test /= 255  
print(x_train.shape[0], 'train samples')  
print(x_test.shape[0], 'test samples')
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-  
datasets/mnist.npz  
11493376/11490434 [=====] - 1s 0us/step  
60000 train samples  
10000 test samples
```

```
[5]: # convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
[6]: model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer=RMSprop(),
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1,
                    validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 512)	262656
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

Total params: 669,706

Trainable params: 669,706

Non-trainable params: 0

Epoch 1/20

469/469 [=====] - 6s 11ms/step - loss: 0.4330 - accuracy: 0.8617 - val_loss: 0.1243 - val_accuracy: 0.9604

Epoch 2/20
469/469 [=====] - 4s 9ms/step - loss: 0.1082 - accuracy: 0.9676 - val_loss: 0.0817 - val_accuracy: 0.9735

Epoch 3/20
469/469 [=====] - 4s 9ms/step - loss: 0.0711 - accuracy: 0.9791 - val_loss: 0.0748 - val_accuracy: 0.9773

Epoch 4/20
469/469 [=====] - 4s 9ms/step - loss: 0.0614 - accuracy: 0.9821 - val_loss: 0.0766 - val_accuracy: 0.9798

Epoch 5/20
469/469 [=====] - 4s 9ms/step - loss: 0.0476 - accuracy: 0.9856 - val_loss: 0.0806 - val_accuracy: 0.9784

Epoch 6/20
469/469 [=====] - 4s 9ms/step - loss: 0.0406 - accuracy: 0.9876 - val_loss: 0.0774 - val_accuracy: 0.9812

Epoch 7/20
469/469 [=====] - 4s 9ms/step - loss: 0.0363 - accuracy: 0.9889 - val_loss: 0.0695 - val_accuracy: 0.9834

Epoch 8/20
469/469 [=====] - 4s 9ms/step - loss: 0.0333 - accuracy: 0.9900 - val_loss: 0.0881 - val_accuracy: 0.9804

Epoch 9/20
469/469 [=====] - 4s 9ms/step - loss: 0.0301 - accuracy: 0.9910 - val_loss: 0.0881 - val_accuracy: 0.9803

Epoch 10/20
469/469 [=====] - 4s 9ms/step - loss: 0.0239 - accuracy: 0.9924 - val_loss: 0.0853 - val_accuracy: 0.9823

Epoch 11/20
469/469 [=====] - 4s 9ms/step - loss: 0.0260 - accuracy: 0.9925 - val_loss: 0.0927 - val_accuracy: 0.9832

Epoch 12/20
469/469 [=====] - 4s 9ms/step - loss: 0.0229 - accuracy: 0.9932 - val_loss: 0.0994 - val_accuracy: 0.9820

Epoch 13/20
469/469 [=====] - 4s 9ms/step - loss: 0.0230 - accuracy: 0.9931 - val_loss: 0.1041 - val_accuracy: 0.9820

Epoch 14/20
469/469 [=====] - 4s 9ms/step - loss: 0.0208 - accuracy: 0.9939 - val_loss: 0.0929 - val_accuracy: 0.9825

Epoch 15/20
469/469 [=====] - 4s 9ms/step - loss: 0.0188 - accuracy: 0.9943 - val_loss: 0.0951 - val_accuracy: 0.9825

Epoch 16/20
469/469 [=====] - 4s 9ms/step - loss: 0.0166 - accuracy: 0.9949 - val_loss: 0.1096 - val_accuracy: 0.9831

Epoch 17/20
469/469 [=====] - 4s 9ms/step - loss: 0.0185 - accuracy: 0.9952 - val_loss: 0.0959 - val_accuracy: 0.9844

```
Epoch 18/20
469/469 [=====] - 4s 9ms/step - loss: 0.0168 -
accuracy: 0.9955 - val_loss: 0.1171 - val_accuracy: 0.9828
Epoch 19/20
469/469 [=====] - 4s 9ms/step - loss: 0.0162 -
accuracy: 0.9952 - val_loss: 0.1145 - val_accuracy: 0.9827
Epoch 20/20
469/469 [=====] - 4s 9ms/step - loss: 0.0155 -
accuracy: 0.9956 - val_loss: 0.1184 - val_accuracy: 0.9833
Test loss: 0.11838609725236893
Test accuracy: 0.983299970626831
```

[]: