

## ex2

June 13, 2021

```
[2]: import sys
from sys import argv
from random import random
from operator import add

from pyspark.sql import SparkSession
```

```
[3]: if __name__ == "__main__":
    """
        Usage: pi [partitions]
    """
    spark = SparkSession\
        .builder\
        .appName("PythonPi")\
        .getOrCreate()

    partitions = abs(int(sys.argv[1], 16)) if len(sys.argv) > 1 else 2
    # partitions = 2
    n = 100000 * partitions

    def f():
        x = random() * 2 - 1
        y = random() * 2 - 1
    #     return 1 if x ** 2 + y ** 2 <= 1 else 0
        if (x ** 2 + y ** 2) <= 1:
            return 1
        else:
            return 0

    count = spark.sparkContext.parallelize(range(1, n + 1), partitions).map(f).
    ↪ reduce(add)
    print("Pi is roughly %f" % (4.0 * count / n))

    spark.stop()
```

Py4JJavaError

Traceback (most recent call last)

```

<ipython-input-3-51f4d32ae6dc> in <module>
    21         return 0
    22
--> 23     count = spark.sparkContext.parallelize(range(1, n + 1), partitions)
    ↪map(f).reduce(add)
    24     print("Pi is roughly %f" % (4.0 * count / n))
    25

/opt/conda/lib/python3.8/site-packages/pyspark/rdd.py in reduce(self, f)
    995         yield reduce(f, iterator, initial)
    996
--> 997         vals = self.mapPartitions(func).collect()
    998         if vals:
    999             return reduce(f, vals)

/opt/conda/lib/python3.8/site-packages/pyspark/rdd.py in collect(self)
    947         """
    948         with SCSiteSync(self.context) as css:
--> 949             sock_info = self.ctx._jvm.PythonRDD.collectAndServe(self.
    ↪_jrdd.rdd())
    950             return list(_load_from_socket(sock_info, self.
    ↪_jrdd_deserializer))
    951

/opt/conda/lib/python3.8/site-packages/py4j/java_gateway.py in __call__(self,
    ↪*args)
    1302
    1303         answer = self.gateway_client.send_command(command)
-> 1304         return_value = get_return_value(
    1305             answer, self.gateway_client, self.target_id, self.name)
    1306

/opt/conda/lib/python3.8/site-packages/pyspark/sql/utils.py in deco(*a, **kw)
    109     def deco(*a, **kw):
    110         try:
--> 111             return f(*a, **kw)
    112         except py4j.protocol.Py4JJavaError as e:
    113             converted = convert_exception(e.java_exception)

/opt/conda/lib/python3.8/site-packages/py4j/protocol.py in
    ↪get_return_value(answer, gateway_client, target_id, name)
    324         value = OUTPUT_CONVERTER[type](answer[2:], gateway_client)
    325         if answer[1] == REFERENCE_TYPE:
-> 326             raise Py4JJavaError(
    327                 "An error occurred while calling {0}{1}{2}.\n".
    328                 format(target_id, ".", name), value)

```

```

Py4JJavaError: An error occurred while calling z:org.apache.spark.api.python.
↳ PythonRDD.collectAndServe.
: org.apache.spark.SparkException: Job aborted due to stage failure: Task 0 in
↳ stage 0.0 failed 1 times, most recent failure: Lost task 0.0 in stage 0.0 (Tl
↳ 0) (jupyter-ntiana55 executor driver): org.apache.spark.api.python.
↳ PythonException: Traceback (most recent call last):
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/worker.py", line 604, i
↳ main
    process()
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/worker.py", line 596, i
↳ process
    serializer.dump_stream(out_iter, outfile)
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/serializers.py", line
↳ 259, in dump_stream
    vs = list(itertools.islice(iterator, batch))
  File "/opt/conda/lib/python3.8/site-packages/pyspark/rdd.py", line 992, in fu c
    initial = next(iterator)
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/util.py", line 73, in
↳ wrapper
    return f(*args, **kwargs)
TypeError: f() takes 0 positional arguments but 1 was given

```

```

    at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.
↳ handlePythonException(PythonRunner.scala:517)
    at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRunner.
↳ scala:652)
    at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRunner.
↳ scala:635)
    at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.
↳ hasNext(PythonRunner.scala:470)
    at org.apache.spark.InterruptibleIterator.hasNext(InterruptibleIterator
↳ scala:37)
    at scala.collection.Iterator.foreach(Iterator.scala:941)
    at scala.collection.Iterator.foreach$(Iterator.scala:941)
    at org.apache.spark.InterruptibleIterator.foreach(InterruptibleIterator
↳ scala:28)
    at scala.collection.generic.Growable.$plus$plus$eq(Growable.scala:62)
    at scala.collection.generic.Growable.$plus$plus$eq$(Growable.scala:53)
    at scala.collection.mutable.ArrayBuffer.$plus$plus$eq(ArrayBuffer.scala
↳ 105)
    at scala.collection.mutable.ArrayBuffer.$plus$plus$eq(ArrayBuffer.scala
↳ 49)
    at scala.collection.TraversableOnce.to(TraversableOnce.scala:315)
    at scala.collection.TraversableOnce.to$(TraversableOnce.scala:313)
    at org.apache.spark.InterruptibleIterator.to(InterruptibleIterator.scal
↳ 28)
    at scala.collection.TraversableOnce.toBuffer(TraversableOnce.scala:307)
    at scala.collection.TraversableOnce.toBuffer$(TraversableOnce.scala:307

```

```

    at org.apache.spark.InterruptibleIterator.toBuffer(InterruptibleIterato
↪ scala:28)
    at scala.collection.TraversableOnce.toArray(TraversableOnce.scala:294)
    at scala.collection.TraversableOnce.toArray$(TraversableOnce.scala:288)
    at org.apache.spark.InterruptibleIterator.toArray(InterruptibleIterator
↪ scala:28)
    at org.apache.spark.rdd.RDD.$anonfun$collect$2(RDD.scala:1030)
    at org.apache.spark.SparkContext.$anonfun$runJob$5(SparkContext.scala:
↪ 2242)
    at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:90)
    at org.apache.spark.scheduler.Task.run(Task.scala:131)
    at org.apache.spark.executor.Executor$TaskRunner.$anonfun$run$3(Executo
↪ scala:497)
    at org.apache.spark.util.Utils$.tryWithSafeFinally(Utils.scala:1439)
    at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:500)
    at java.base/java.util.concurrent.ThreadPoolExecutor.
↪ runWorker(ThreadPoolExecutor.java:1128)
    at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.
↪ run(ThreadPoolExecutor.java:628)
    at java.base/java.lang.Thread.run(Thread.java:834)

```

#### Driver stacktrace:

```

    at org.apache.spark.scheduler.DAGScheduler.
↪ failJobAndIndependentStages(DAGScheduler.scala:2253)
    at org.apache.spark.scheduler.DAGScheduler.
↪ $anonfun$abortStage$2(DAGScheduler.scala:2202)
    at org.apache.spark.scheduler.DAGScheduler.
↪ $anonfun$abortStage$2$adapted(DAGScheduler.scala:2201)
    at scala.collection.mutable.ResizableArray.foreach(ResizableArray.scala
↪ 62)
    at scala.collection.mutable.ResizableArray.foreach$(ResizableArray.scal
↪ 55)
    at scala.collection.mutable.ArrayBuffer.foreach(ArrayBuffer.scala:49)
    at org.apache.spark.scheduler.DAGScheduler.abortStage(DAGScheduler.scal
↪ 2201)
    at org.apache.spark.scheduler.DAGScheduler.
↪ $anonfun$handleTaskSetFailed$1(DAGScheduler.scala:1078)
    at org.apache.spark.scheduler.DAGScheduler.
↪ $anonfun$handleTaskSetFailed$1$adapted(DAGScheduler.scala:1078)
    at scala.Option.foreach(Option.scala:407)
    at org.apache.spark.scheduler.DAGScheduler.
↪ handleTaskSetFailed(DAGScheduler.scala:1078)
    at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.
↪ doOnReceive(DAGScheduler.scala:2440)
    at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.
↪ onReceive(DAGScheduler.scala:2382)
    at org.apache.spark.scheduler.DAGSchedulerEventProcessLoop.
↪ onReceive(DAGScheduler.scala:2371)

```

```

at org.apache.spark.util.EventLoop$$anon$1.run(EventLoop.scala:49)
at org.apache.spark.scheduler.DAGScheduler.runJob(DAGScheduler.scala:86)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2202)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2223)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2242)
at org.apache.spark.SparkContext.runJob(SparkContext.scala:2267)
at org.apache.spark.rdd.RDD.$anonfun$collect$1(RDD.scala:1030)
at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.
↳ scala:151)
    at org.apache.spark.rdd.RDDOperationScope$.withScope(RDDOperationScope.
↳ scala:112)
        at org.apache.spark.rdd.RDD.withScope(RDD.scala:414)
        at org.apache.spark.rdd.RDD.collect(RDD.scala:1029)
        at org.apache.spark.api.python.PythonRDD$.collectAndServe(PythonRDD.
↳ scala:180)
            at org.apache.spark.api.python.PythonRDD.collectAndServe(PythonRDD.scal
            at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.
↳ invoke0(Native Method)
                at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.
↳ invoke(NativeMethodAccessorImpl.java:62)
                    at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.
↳ invoke(DelegatingMethodAccessorImpl.java:43)
                        at java.base/java.lang.reflect.Method.invoke(Method.java:566)
                        at py4j.reflection.MethodInvoker.invoke(MethodInvoker.java:244)
                        at py4j.reflection.ReflectionEngine.invoke(ReflectionEngine.java:357)
                        at py4j.Gateway.invoke(Gateway.java:282)
                        at py4j.commands.AbstractCommand.invokeMethod(AbstractCommand.java:132)
                        at py4j.commands.CallCommand.execute(CallCommand.java:79)
                        at py4j.GatewayConnection.run(GatewayConnection.java:238)
                        at java.base/java.lang.Thread.run(Thread.java:834)
Caused by: org.apache.spark.api.python.PythonException: Traceback (most recent
↳ call last):
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/worker.py", line 604, i
↳ main
    process()
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/worker.py", line 596, i
↳ process
    serializer.dump_stream(out_iter, outfile)
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/serializers.py", line
↳ 259, in dump_stream
    vs = list(itertools.islice(iterator, batch))
  File "/opt/conda/lib/python3.8/site-packages/pyspark/rdd.py", line 992, in fu c
    initial = next(iterator)
  File "/usr/local/spark/python/lib/pyspark.zip/pyspark/util.py", line 73, in
↳ wrapper
    return f(*args, **kwargs)
TypeError: f() takes 0 positional arguments but 1 was given

```

```

    at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.
↪handlePythonException(PythonRunner.scala:517)
    at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRunner.
↪scala:652)
    at org.apache.spark.api.python.PythonRunner$$anon$3.read(PythonRunner.
↪scala:635)
    at org.apache.spark.api.python.BasePythonRunner$ReaderIterator.
↪hasNext(PythonRunner.scala:470)
    at org.apache.spark.InterruptibleIterator.hasNext(InterruptibleIterator
↪scala:37)
    at scala.collection.Iterator.foreach(Iterator.scala:941)
    at scala.collection.Iterator.foreach$(Iterator.scala:941)
    at org.apache.spark.InterruptibleIterator.foreach(InterruptibleIterator
↪scala:28)
    at scala.collection.generic.Growable.$plus$plus$eq(Growable.scala:62)
    at scala.collection.generic.Growable.$plus$plus$eq$(Growable.scala:53)
    at scala.collection.mutable.ArrayBuffer.$plus$plus$eq(ArrayBuffer.scala
↪105)
    at scala.collection.mutable.ArrayBuffer.$plus$plus$eq(ArrayBuffer.scala
↪49)
    at scala.collection.TraversableOnce.to(TraversableOnce.scala:315)
    at scala.collection.TraversableOnce.to$(TraversableOnce.scala:313)
    at org.apache.spark.InterruptibleIterator.to(InterruptibleIterator.scal. :
↪28)
    at scala.collection.TraversableOnce.toBuffer(TraversableOnce.scala:307)
    at scala.collection.TraversableOnce.toBuffer$(TraversableOnce.scala:307)
    at org.apache.spark.InterruptibleIterator.toBuffer(InterruptibleIterato.
↪scala:28)
    at scala.collection.TraversableOnce.toArray(TraversableOnce.scala:294)
    at scala.collection.TraversableOnce.toArray$(TraversableOnce.scala:288)
    at org.apache.spark.InterruptibleIterator.toArray(InterruptibleIterator
↪scala:28)
    at org.apache.spark.rdd.RDD.$anonfun$collect$2(RDD.scala:1030)
    at org.apache.spark.SparkContext.$anonfun$runJob$5(SparkContext.scala:
↪2242)
    at org.apache.spark.scheduler.ResultTask.runTask(ResultTask.scala:90)
    at org.apache.spark.scheduler.Task.run(Task.scala:131)
    at org.apache.spark.executor.Executor$TaskRunner.$anonfun$run$3(Executo.
↪scala:497)
    at org.apache.spark.util.Utils$.tryWithSafeFinally(Utils.scala:1439)
    at org.apache.spark.executor.Executor$TaskRunner.run(Executor.scala:500)
    at java.base/java.util.concurrent.ThreadPoolExecutor.
↪runWorker(ThreadPoolExecutor.java:1128)
    at java.base/java.util.concurrent.ThreadPoolExecutor$Worker.
↪run(ThreadPoolExecutor.java:628)
    ... 1 more

```

[ ]: